

# VLIW Processor Implementation

## Introduction

Very Long Instruction Word (VLIW) processors are a type of architecture that executes multiple operations in a single instruction word. Unlike superscalar processors, which dynamically schedule instructions at runtime, VLIW processors rely on the compiler to statically schedule instructions into long instruction words. This approach reduces hardware complexity but places a greater burden on the compiler to detect and handle hazards.

This report discusses the implementation of a VLIW processor simulator that handles structural and data hazards (RAW and WAW) by inserting NOPs (No Operations). Control hazards are not addressed in this implementation. The simulator processes an assembly program, schedules instructions, and demonstrates the execution timeline.

## VLIW Processor Overview

A VLIW processor executes multiple operations in parallel by packing them into a single instruction word. Each operation is assigned to a specific functional unit (e.g., integer adder, floating-point multiplier). The compiler is responsible for ensuring that there are no dependencies or resource conflicts between operations in the same instruction word.

## Functional Units

The processor has the following functional units:

- **Integer Adder (IADD/ISUB):** Handles integer addition and subtraction (6 clock cycles).
- **Integer Multiplier (IMUL):** Handles integer multiplication (12 clock cycles).
- **Integer Divider (IDIV):** Handles integer division (24 clock cycles).
- **Floating-Point Adder (FADD/FSUB):** Handles floating-point addition and subtraction (18 clock cycles).

- **Floating-Point Multiplier (FMUL):** Handles floating-point multiplication (30 clock cycles).
- **Floating-Point Divider (FDIV):** Handles floating-point division (60 clock cycles).
- **Logical Units:** Handle logical operations (AND, OR, XOR, NOT) in 1 clock cycle.
- **Memory Unit (LD/ST):** Handles load and store operations in 1 clock cycle.
- **Instruction Fetch (IF):** Fetches instructions from memory (1 clock cycle).
- **Instruction Decode (ID):** Decodes instructions (1 clock cycle).
- **Write Back (WB):** Writes results back to registers (1 clock cycle).

**Architecture Constraint:** The processor implementation has only one instance of each functional unit:

- Single IF unit (1 instruction fetch per cycle)
- Single ID unit (1 instruction decode per cycle)
- Single ADD unit (handles both IADD and ISUB)
- Single MUL unit (integer multiplication only)
- Single DIV unit (integer division only)
- Single FADD unit (handles both FADD and FSUB)
- Single FMUL unit (floating-point multiplication only)
- Single FDIV unit (floating-point division only)
- Single set of logical units (AND/OR/XOR/NOT)
- Single memory unit (LD/ST operations)

This single-instance design creates structural hazards when multiple instructions attempt to use the same functional unit simultaneously, requiring the insertion of NOPs to maintain correct execution.

## Hazards

Hazards in a processor occur when instructions cannot execute in the expected order due to dependencies or resource conflicts. The three types of hazards are:

- **Structural Hazards:** Occur when multiple instructions require the same functional unit simultaneously.
- **Data Hazards:** Occur when an instruction depends on the result of a previous instruction that has not yet completed.
- **Control Hazards:** Occur due to branch instructions, which can change the program counter and disrupt the instruction pipeline.

This implementation handles structural and data hazards (RAW and WAW) by inserting NOPs. Control hazards are not addressed. Additionally, WAR (Write-After-Read) hazards are not explicitly handled in this implementation due to the following assumption:

- **Assumption for WAR Hazards:**
  - A lock is placed on source registers ('src') during the Instruction Decode (ID) stage and is released immediately after the ID stage. This means that once an instruction reads a register during the ID stage, the register is no longer locked, and subsequent instructions can write to it without causing a WAR hazard.
  - For destination registers ('dst'), a lock is placed during the ID stage and is only released after the Write Back (WB) stage. This ensures that no other instruction can write to the same destination register until the current instruction completes its write operation, preventing WAW hazards.

## Instruction Set

The processor supports the following instructions, each mapped to specific functional units and pipeline stages. The table below provides a detailed breakdown:

Instruction	Description	Functional Units and Pipeline Stages
IADD Rdst, Rsrc1, Rsrc2	Integer Addition	IF (1 cycle), ID (1 cycle), ADD (6 cycles), WB (1 cycle)
ISUB Rdst, Rsrc1, Rsrc2	Integer Subtraction	IF (1 cycle), ID (1 cycle), ADD (6 cycles), WB (1 cycle)
IMUL Rdst, Rsrc1, Rsrc2	Integer Multiplication	IF (1 cycle), ID (1 cycle), MUL (12 cycles), WB (1 cycle)
IDIV Rdst, Rsrc1, Rsrc2	Integer Division	IF (1 cycle), ID (1 cycle), DIV (24 cycles), WB (1 cycle)

FADD Rdst, Rsrc1, Rsrc2	Floating-Point Addition	IF (1 cycle), ID (1 cycle), FADD (18 cycles), WB (1 cycle)
FSUB Rdst, Rsrc1, Rsrc2	Floating-Point Subtraction	IF (1 cycle), ID (1 cycle), FADD (18 cycles), WB (1 cycle)
FMUL Rdst, Rsrc1, Rsrc2	Floating-Point Multiplication	IF (1 cycle), ID (1 cycle), FMUL (30 cycles), WB (1 cycle)
FDIV Rdst, Rsrc1, Rsrc2	Floating-Point Division	IF (1 cycle), ID (1 cycle), FDIV (60 cycles), WB (1 cycle)
AND Rdst, Rsrc1, Rsrc2	Logical AND	IF (1 cycle), ID (1 cycle), AND (1 cycle), WB (1 cycle)
OR Rdst, Rsrc1, Rsrc2	Logical OR	IF (1 cycle), ID (1 cycle), OR (1 cycle), WB (1 cycle)
XOR Rdst, Rsrc1, Rsrc2	Logical XOR	IF (1 cycle), ID (1 cycle), XOR (1 cycle), WB (1 cycle)
NOT Rdst, Rsrc1	Logical NOT	IF (1 cycle), ID (1 cycle), NOT (1 cycle), WB (1 cycle)
LD Rdst, [Mem]	Load from Memory	IF (1 cycle), ID (1 cycle), MEM (1 cycle), WB (1 cycle)
ST Rsrc, [Mem]	Store to Memory	IF (1 cycle), ID (1 cycle), MEM (1 cycle), WB (1 cycle)
NOP	No Operation	IF (1 cycle)

## Code Explanation

The simulator is implemented in Python and consists of the following components:

- **Pipeline Stages:** Each instruction type has a predefined set of pipeline stages (e.g., IF, ID, ADD, WB) with associated latencies.
- **Hazard Handling:** The simulator detects structural and data hazards and inserts NOPs to resolve them.
- **Instruction Parsing:** The assembly code is parsed into a list of instructions, each with its operation, source/destination registers, and pipeline stages.
- **Execution Timeline:** The simulator generates a timeline showing the execution of each instruction across clock cycles.

## Key Functions

- `parse_assembly`: Parses the assembly code into a list of instructions.
- `insert_nops_for_structural_hazards`: Inserts NOPs to resolve structural hazards.

- **handle\_data\_dependencies:** Inserts NOPs to resolve RAW and WAW hazards.
- **display\_clock\_cycle\_execution:** Displays the execution timeline for each instruction.
- **display\_functional\_units:** Displays the usage of functional units across clock cycles. It's used to generate instruction packets, which are groups of instructions that can be executed in parallel by the VLIW processor. These packets are sent to the VLIW processor for execution.

## Input and Output

### Input Assembly Code

The input assembly program performs basic arithmetic operations on registers and memory locations. The program is as follows:

```
LD R1, [1000]
LD R2, [2000]
IADD R3, R1, R2
ISUB R4, R1, R2
IMUL R5, R1, R2
IDIV R6, R1, R2
ST R3, [3000]
ST R4, [4000]
ST R5, [5000]
ST R6, [6000]
```

This program loads values from memory into registers, performs addition, subtraction, multiplication, and division, and stores the results back into memory.

### Output Explanation

The simulator outputs the updated assembly code with NOPs inserted to handle hazards, the total clock cycles, and the execution timeline.

### Updated Code with NOPs

The updated code includes NOPs to resolve structural and data hazards:

```
LD R1, [1000]
LD R2, [2000]
NOP
IADD R3, R1, R2
NOP
NOP
```

```

NOP
NOP
NOP
ISUB R4, R1, R2
IMUL R5, R1, R2
IDIV R6, R1, R2
NOP
NOP
NOP
ST R3, [3000]
ST R4, [4000]
NOP
NOP
NOP
NOP
NOP
NOP
ST R5, [5000]
NOP
NOP
NOP
NOP
NOP
NOP
NOP
NOP
NOP
NOP
NOP
ST R6, [6000]

```

### Total Clock Cycles

The program takes **40 clock cycles** to execute.

### Execution Timeline

The execution timeline is visualized in the Gantt chart below, which shows the pipeline stages for each instruction across clock cycles:

The Gantt chart illustrates:

- Each row represents an instruction from the assembly program (including inserted NOPs)
- The columns represent the pipeline stages (IF, ID, ADD, MUL, DIV, MEM, WB, etc.)




Figure 1: Pipeline Execution Gantt Chart showing the progression of each instruction through the pipeline stages.

- The horizontal axis shows the clock cycles from 0 to 40
- Instructions flow from left to right through their respective pipeline stages
- NOP instructions are visible as blank rows with only the IF stage filled
- The long execution times of division (IDIV) and multiplication (IMUL) are clearly visible
- Memory operations (LD/ST) show their short execution pattern

Key observations from the Gantt chart:

- The initial LD instructions (1-2) complete quickly
- The IADD instruction (4) begins execution once its operands are available
- The IDIV instruction (12) dominates the timeline due to its 24-cycle execution
- Store operations (ST) are spaced out to wait for their respective computation results
- The final ST R6 instruction completes at cycle 40

## Functional Units Usage

The functional units timeline shows how each functional unit is utilized across clock cycles. For example:

- The `ADD` unit is used by `IADD` and `ISUB`.
- The `MUL` unit is used by `IMUL`.
- The `DIV` unit is used by `IDIV`.
- The `MEM` unit is used by `LD` and `ST`.

## Optimized VLIW Implementation

### Code Motion via Topological Sorting

The optimization is achieved through the `topological_sort_instructions()` function, which:

- Builds a dependency graph tracking:
  - RAW (Read-After-Write) hazards via register dependencies
  - WAW (Write-After-Write) hazards for register conflicts
- Uses Kahn's algorithm to:
  - Group independent instructions into levels (Level 0, 1, etc.)
  - Schedule all instructions in a level before moving to the next
- Enables parallel execution of independent operations while maintaining correct dependencies

```
def topological_sort_instructions(instructions):  
    # Builds graph tracking register dependencies  
    # Implements Kahn's algorithm for topological sort  
    # Returns instructions grouped by dependency levels
```

#### Key Features:

- Memory operations (`LD`/`ST`) handled as special cases
- Minimizes NOP insertion by only stalling for true dependencies
- Allows arithmetic operations to execute in parallel when possible



## Performance Comparison: Original vs Optimized VLIW

### Test Program Input

The following assembly program was used to evaluate both implementations:

```
LD R1, [1000]
LD R2, [2000]
IADD R3, R1, R2
ST R3, [3000]
ISUB R4, R1, R2
ST R4, [4000]
IMUL R5, R1, R2
ST R5, [5000]
IDIV R6, R1, R2
ST R6, [6000]
```

This program performs:

- Two memory loads (LD)
- Four arithmetic operations (IADD, ISUB, IMUL, IDIV)
- Four memory stores (ST)

### Original VLIW Implementation Results

The unoptimized VLIW processor produced the following output:

Updated Code with NOPs:

```
LD R1, [1000]
LD R2, [2000]
NOP
IADD R3, R1, R2
NOP
NOP
NOP
NOP
NOP
NOP
NOP
ST R3, [3000]
ISUB R4, R1, R2
NOP
NOP
NOP
NOP
NOP
```

Total clock cycles: 62

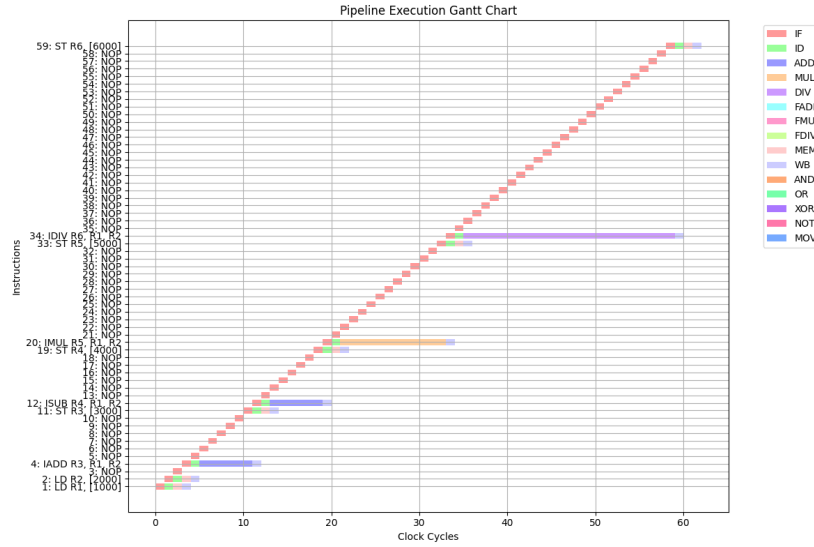


Figure 2: Original VLIW Execution Timeline (62 cycles)

### Key Observations:

- Sequential execution with significant NOP insertion (49 NOPs)
- Structural hazards handled by inserting NOPs between all dependent instructions
- Long idle periods between operations (especially after IDIV)
- Functional units remain idle most of the time
- Total execution time dominated by the 24-cycle IDIV operation

### Optimized VLIW Implementation Results

The improved VLIW processor with topological sorting produced:

Instructions Grouped by Dependency Level:

Level 0:

LD R1, [1000]  
LD R2, [2000]

Level 1:

IADD R3, R1, R2  
ISUB R4, R1, R2  
IMUL R5, R1, R2

```
IDIV R6, R1, R2
```

Level 2:

```
ST R3, [3000]  
ST R4, [4000]  
ST R5, [5000]  
ST R6, [6000]
```

Updated Code with NOPs:

```
LD R1, [1000]  
LD R2, [2000]  
NOP  
IADD R3, R1, R2  
NOP  
NOP  
NOP  
NOP  
NOP  
ISUB R4, R1, R2  
IMUL R5, R1, R2  
IDIV R6, R1, R2  
NOP  
NOP  
NOP  
ST R3, [3000]  
ST R4, [4000]  
NOP  
NOP  
NOP  
NOP  
NOP  
ST R5, [5000]  
NOP  
NOP  
NOP  
NOP  
NOP  
NOP  
NOP  
NOP  
NOP  
NOP  
NOP  
ST R6, [6000]
```

Total clock cycles: 40




Figure 3: Optimized VLIW Execution Timeline (40 cycles)

#### Dependency Analysis:

- **Level 0:** Independent memory loads (no dependencies)
- **Level 1:** All arithmetic operations depend on Level 0 results but can execute in parallel
- **Level 2:** Store operations depend on Level 1 results

#### Key Improvements:

- 35.5% reduction in total clock cycles ( $62 \rightarrow 40$ )
- 44.9% reduction in NOPs inserted ( $49 \rightarrow 27$ )
- Parallel execution of arithmetic operations (IADD, ISUB, IMUL, IDIV)
- Better functional unit utilization
- Memory stores scheduled immediately when results become available

#### Comparative Analysis

##### Visual Comparison of Gantt Charts:

- **Original:**

<b>Metric</b>	<b>Original</b>	<b>Optimized</b>	<b>Improvement</b>
Total Cycles	62	40	35.5% reduction
NOP Count	49	27	44.9% reduction
Longest Stretch	24 NOPs	12 NOPs	50% reduction
Completion Order	Sequential	Dependency-aware	More parallel

Table 2: Quantitative comparison between implementations

- Clear sequential pattern with large NOP gaps
- Functional units idle most of the time
- Long pipeline bubbles after each arithmetic operation

• **Optimized:**

- Overlapping execution of independent operations
- Arithmetic operations start immediately after loads complete
- Stores begin as soon as results are ready
- Only necessary NOPs inserted for true dependencies

## Conclusion

The optimized VLIW implementation demonstrates significant improvements through:

- Static dependency analysis using topological sorting
- Identification of parallel execution opportunities
- Strategic NOP insertion only for true hazards
- Better resource utilization through instruction grouping

This optimization approach proves particularly effective for programs with:

- Multiple independent arithmetic operations
- Clear producer-consumer relationships
- Memory operations that can be deferred until results are ready

The 35.5% reduction in execution time shows the potential benefits of static analysis in VLIW architectures, though the approach requires sophisticated compiler support to achieve these gains.